

CS 228: Computational Learning Theory
Homework 2
Due: March 12, 11:59PM

Policy reminders: You are strongly encouraged to type your solutions using LATEX. You may discuss problems with your classmates, but not merely copy each others solutions. You must write all solutions by yourself, list your collaborators on your problem sets and also appropriately cite any resources outside of the class materials that you have used. You are not allowed to look up solutions to the problems. Please do not use LLMs or LLM-assisted tools for finding solutions to the problems.

Problem 1 (10pts). Efficient Evaluability is Necessary. The definition of learning in the PAC model requires that hypotheses produced by a learning algorithms are evaluable efficiently (that is, in polynomial time). In this problem you will show that this requirement is necessary for a meaningful notion of learning.

Suppose we consider relaxing this restriction, and let \mathcal{H} be the class of all Turing machines (not necessarily polynomial time). Show that if \mathcal{C}_n is the class of all Boolean circuits of size at most $p(n)$ for some fixed polynomial $p(\cdot)$, then \mathcal{C} is efficiently PAC learnable using \mathcal{H} .

Argue that your solution shows that this relaxation trivializes the model of learning. (Hint: How could you use your solution to trivially learn other concept classes we've seen in class?)

To show that the requirement that the hypotheses produced by a learning algorithm are efficiently evaluable in polynomial time, consider the following scenario: Suppose we have a learning algorithm \mathcal{L} that efficiently PAC learns the concept class \mathcal{C}_n of Boolean circuits with size bounded by $p(n)$:

Algorithm \mathcal{L} :

- **Input:** A sample $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ where each $y_i = c(x_i)$ for some target concept $c \in \mathcal{C}_n$
- **Output:** A Turing machine T_S that implements a hypothesis h

Definition of T_S : For any input $x \in \{0, 1\}^n$, T_S executes:

1. Systematically enumerate all Boolean circuits $c' \in \mathcal{C}_n$ in order of increasing circuit size
2. For each circuit c' , check if $c'(x_i) = y_i$ for all $(x_i, y_i) \in S$
3. Return $c'(x)$ for the first circuit c' is found that is consistent with S

Proof \mathcal{L} is an Occam Algorithm: Since the target concept $c \in \mathcal{C}_n$ is consistent with S , the algorithm will find at least one consistent circuit (possibly c itself or another circuit that agrees with c on all examples in S). Also, the description size of T_S is polynomial in n and m . Specifically, the Turing machine T_S can be described in $O(n + \log m)$ bits as it only has to encode the enumeration procedure for circuits (constant-sized code), the sample S (requires at most $O(m \cdot n)$ bits), and the consistency checking logic (constant-sized code). Thus, for some constants α and β , we have $\text{size}(T_S) \leq (n \cdot \text{size}(c))^\alpha \cdot m^\beta$. Lastly, constructing T_S takes polynomial time since we only need to encode the sample and search procedure. Therefore, \mathcal{L} is an Occam algorithm that efficiently PAC learns \mathcal{C}_n using \mathcal{H} as the hypothesis class.

Conclusion: This approach can be generalized to learn *any* concept class that has a finite representation. For any learnable concept class \mathcal{D} , we can define a similar Turing machine that enumerates all concepts in \mathcal{D} and returns the first one consistent with the sample. Despite such a Turing machine having a compact description making the training process efficient, evaluating it on new examples would be computationally expensive. The Turing machine could have to do an exhaustive search through an exponentially large search space of possible concepts before producing an output. Thus, without the efficiently evaluable constraint, we could trivially learn any concept class by outputting a hypothesis that performs a brute-force search through the concept space—a strategy that doesn't reflect the true nature of learning.

Problem 2 (15pts). Infinite Mistake-Bounded Model. The *Infinite Mistake-Bounded (IMB) model* is a Boolean on-line learning model where the number of attributes in the world may be infinite. It is assumed, however, that each example has a finite number of attributes assigned 1. An example is presented to the learner as a list of its positive attributes. The learning scenario is the same as in the (standard) mistake-bounded model.

Provide an on-line learning algorithm that makes at most $\mathcal{O}(k \log n)$ mistakes when learning any monotone disjunction of k literals, if every example presented to the learner has at most n positive attributes.

Note: Assume for full credit that the learner does not know the value of n in advance. For partial credit, you can use the value of n in your algorithm

Problem 3 (15 points). Function Classes: Linear Thresholds, DNFs, and Decision Lists.

- Prove that any 1-decision list over x_1, \dots, x_n can be expressed as a linear threshold function $\mathbb{1}[w \cdot x \geq \theta]$, where $w \in \mathbb{R}^n \setminus \{0\}$, $\theta \in \mathbb{R}$.
- Prove that the class of linear threshold functions and the class of $\text{poly}(n)$ -term DNF are incomparable in that neither class is contained in the other.

Problem 4 (15 points). Perceptrons Can Make Exponentially Many Mistakes.

- The *margin* δ of a set of points $X \subseteq \{0, 1\}^n$ labeled by a function f is defined as follows.

$$\delta = \max_{w \in \mathbb{R}^n, \|w\|_2=1, \theta \in \mathbb{R}} \{\delta' \mid \forall x \in X, |w \cdot x - \theta| \geq \delta' \text{ and } (w \cdot x \geq \theta \text{ iff } f(x) = 1)\}$$

(This is the δ of the Perceptron algorithm.) Give a set of $\mathcal{O}(n)$ examples on $\{0, 1\}^n$ that are linearly separable but for which the margin is exponentially small in n .

- Give an example of a linear threshold function on which the Perceptron algorithm can make exponentially many mistakes.

Problem 5 (10pts). Logarithmic Mistake Bounds.

- Let \mathcal{C}_n be a finite concept class. Give a learning algorithm for \mathcal{C}_n which has mistake bound $\log_2 |\mathcal{C}_n|$. Your algorithm need not be computationally efficient.
- Show a concept class \mathcal{C} for which there exists an algorithm with a mistake bound that is asymptotically better than $\log_2 |\mathcal{C}_n|$.

Problem 6 (15 points). Robust Winnow Algorithm. In this question we examine the error robustness of the Winnow algorithm. Let f be a monotone disjunction of k variables. Suppose that an adversary manipulates the labels of some of the examples seen by the Winnow algorithm as it

runs on f : during the online learning process, for some inputs x , the Winnow algorithm is told that the label of x is $1 - f(x)$ when the correct value was actually $f(x)$. When $f(x) = 1$ but the adversary flips its reported label to 0, we call this is a *false negative example*. When $f(x) = 0$ but the algorithm is presented with the label 1, this is called a *false positive example*.

Assume that the algorithm sees s false negative examples and t false positive examples during the course of its learning. Prove that the total number of mistakes the Winnow algorithm¹ will make (on examples that are labeled correctly) is $O(k \log n + ks + t)$.

Note. You will have to modify the Winnow algorithm to get this result. In class the demotion step set all weights $w_i = 0$ if $x_i = 1$. Argue why this will not work in the above setting and show that if instead you set $w_i \leftarrow w_i/2$, you will get the required bounds.

Problem 7 (15pts). VC Dimension of Linear Halfspaces. Let \mathcal{C} be the concept class of linear halfspaces in \mathbb{R}^n . A halfspace is specified by an inequality of the form $c(x) = \mathbb{1}[w \cdot x \geq \theta]$, where $w \in \mathbb{R}^n \setminus \{0\}$ and $\theta \in \mathbb{R}$. Prove that the VC dimension of \mathcal{C} is $n + 1$.

Note. You may use the following result without proof:

Theorem 1 (Radon's theorem). Let $S = \{x^{(1)}, \dots, x^{(m)}\} \subset \mathbb{R}^n$ be a set of m points in \mathbb{R}^n . The convex hull of S is the set

$$\{z \in \mathbb{R}^n \mid \exists \lambda_1, \dots, \lambda_m \text{ s.t. } z = \sum_{i=1}^m \lambda_i x^{(i)}, \text{ with each } \lambda_i \geq 0 \text{ and } \sum_{i=1}^m \lambda_i = 1\}$$

If $m \geq n + 2$ then S must have two disjoint subsets S_1 and S_2 whose convex hulls intersect.

¹if modified appropriately; see the note at the end of the problem statement.